

Localizing Persona Representations in LLMs




Figure 6: (Q1): Euclidean distances stratified by dimensions across models and layers.